

# DSA5105 Principles of Machine Learning

AY2024/25 Sem1 By Zhao Peiduo

## Supervised Learning

### Lecture 1

**Supervised Learning** Supervised learning is the most common type of machine learning problem, where the goal is to make predictions and learn a function that maps an input to an output based on example input-output pairs.

**Problem Setup** Given a dataset  $\mathcal{D} = \{(x_i, y_i)\}_{i=1}^N$ , where  $x_i$  are the inputs and  $y_i$  are the corresponding outputs or labels, the goal is to learn the mapping  $f: X \rightarrow Y$  such that  $f(x_i) \approx y_i$  for all  $i$ .

**Hypothesis Space** The **oracle**  $f^*$  is unknown to us except through the dataset. The function  $f$  is chosen from a hypothesis space  $\mathcal{H}$ , which is a set of candidate functions. For example, in linear regression, the hypothesis space might be  $\mathcal{H} = \{f: f(x) = w_0 + w_1 x \mid w_0, w_1 \in \mathbb{R}\}$ .

### Three Paradigms of Supervised Learning

- **Approximation:** Analyze the breadth and depth of the hypothesis space  $\mathcal{H}$  to determine if it contains, or closely approximates, the optimal function. (How large is our hypothesis space?)
- **Optimization:** Design and implement efficient algorithms to address the empirical risk minimization problem and find or closely approximate the best function within  $\mathcal{H}$ . (How can we find or get close to an approximation  $\hat{f}$  of  $f^*$ ?)
- **Generalization:** Evaluate whether the optimized model can effectively generalize to new, unseen data, focusing on the interplay between data size and model complexity. This is done through minimizing the population risk:  $R_{\text{pop}}(f) = \mathbb{E}_{(x,y) \sim P}[L(y, f(x))]$ . (Can the  $\hat{f}$  found generalized to unseen examples?)

**Empirical Risk Minimization (ERM)** The learning process aims to find a function  $f \in \mathcal{H}$  that minimizes the empirical risk,  $R_{\text{emp}}(f) = \frac{1}{N} \sum_{i=1}^N L(y_i, f(x_i))$ , where  $y_i = f^*(x_i)$ .

**Loss Function** To quantify how well a function  $f$  fits the data, we use a loss function  $L(y, \hat{y})$ , where  $y$  is the true output, and  $\hat{y} = f(x)$  is the predicted output. Common loss functions include the mean squared error (MSE) for regression tasks:  $L(y, \hat{y}) = \frac{1}{2}(y - \hat{y})^2$

**Ordinary Least Squares Formula** The formula for Ordinary Least Squares in a simple linear regression context is given by:  $\beta = (X^T X)^{-1} X^T y$ . In 1D context, the empirical risk minimization problem can be defined as  $\min_{w_0, w_1} \frac{1}{2N} \sum_{i=1}^N (w_0 + w_1 x_i - y_i)^2$ . By setting the partial derivatives to zero, the ordinary least squares formula 1D is given by:  $\hat{w}_0 = \bar{y} - \hat{w}_1 \bar{x}$   $\hat{w}_1 = \frac{\sum_i (x_i - \bar{x})(y_i - \bar{y})}{\sum_i (x_i - \bar{x})^2}$   $\bar{x} = \frac{1}{N} \sum_i x_i$   $\bar{y} = \frac{1}{N} \sum_i y_i$

**Huber Loss** The Huber loss, used for robust regression, is defined as:  $L_\delta(y, \hat{y}) = \begin{cases} \frac{1}{2}(y - \hat{y})^2 & \text{if } |y - \hat{y}| \leq \delta \\ \delta|y - \hat{y}| - \frac{1}{2}\delta^2 & \text{otherwise} \end{cases}$

**General Ordinary Least Squares Formula** Consider  $x \in \mathbb{R}^d$  and the new hypothesis space  $\mathcal{H}_M = \{f: f(x) = \sum_{j=0}^{M-1} w_j \varphi_j(x)\}$  Each  $\varphi_j: \mathbb{R}^d \rightarrow \mathbb{R}$  is called a **basis function** or **feature map**.

We can rewrite the ERM  $\min_{w \in \mathbb{R}^M} \frac{1}{2N} \sum_{i=1}^N \left( \sum_{j=0}^{M-1} w_j \varphi_j(x_i) - y_i \right)^2$  into  $\min_{w \in \mathbb{R}^M} \frac{1}{2N} \|\Phi w - y\|^2$ . Solving by setting  $\nabla R_{\text{emp}}(\hat{w}) = 0$ , we have  $\hat{w} = (\Phi^T \Phi)^{-1} \Phi^T y$ , given invertible  $\Phi^T \Phi$ .

For cases where  $\Phi^T \Phi$  is not invertible, the formula using the Moore-Penrose pseudoinverse is:  $\hat{w}(u) = \Phi^\dagger y + (I - \Phi^\dagger \Phi) u$   $u \in \mathbb{R}^M$

**Overfitting:** Overfitting occurs when the hypothesis space  $\mathcal{H}$  is too large, allowing the model to fit the noise in the training data. This results in poor generalization to new data.

**Regularization:** To prevent overfitting, regularization techniques add a penalty to the loss function:  $\min_{w \in \mathbb{R}^M} \frac{1}{2N} \|\Phi w - y\|^2 + \lambda C(w)$   $\frac{1}{2N} \|\Phi w - y\|^2 + \lambda C(w)$ . Minimizing the ERM we get  $\hat{w} = (\Phi^T \Phi + \lambda I_M)^{-1} \Phi^T y$  which is always invertible for positive  $\lambda$ .

**Common regularization terms:**

- **L<sub>2</sub> (Ridge) regularization:**  $\lambda \sum_{j=1}^p w_j^2$
- **L<sub>1</sub> (Lasso) regularization:**  $\lambda \sum_{j=1}^p |w_j|$

**Softmax Function** For a multi-class classification problem with  $K$  classes, the softmax function is defined as:  $\text{softmax}(z_i) = \frac{\exp(z_i)}{\sum_{j=1}^K \exp(z_j)}$

**Cross-Entropy Loss** Commonly used in classification tasks, the cross-entropy loss is:  $L(y, p) = - \sum_i y_i \log(p_i)$

### Lecture 2

**Another view of feature maps** One can also view feature maps as implicitly defining some sort of **similarity measure**. Consider two vectors  $u$  and  $v$ . Then,  $u^T v$  measures how similar they are. Feature maps define a **similarity** between two samples  $x, x'$  by computing the dot product in **feature space**.

**Reformulation of Ridge Regression** We rewrite the regularized least squares solution in another way:

$$\hat{w} = (\Phi^T \Phi + \lambda I_M)^{-1} \Phi^T y = \Phi^T (\Phi \Phi^T + \lambda I_N)^{-1} y$$

**Proof:**

$$(\Phi^T \Phi + \lambda I_M) \Phi^T = \Phi^T (\Phi \Phi^T + \lambda I_N)^{-1}$$

$$= \Phi^T (\Phi \Phi^T + \lambda I_N)^{-1} y$$

$$= \Phi^T \Phi^T y (\Phi^T + \lambda I_N)^{-1}$$

$$\lambda I_N^{-1} y = (\Phi^T \Phi + \lambda I_M)^{-1} \Phi^T y$$

### Reformulation of Ridge Regression

$$\hat{f}(x) = \sum_{i=1}^N \alpha_i \varphi(x_i)^\top \varphi(x)$$

$$\alpha = (G + \lambda I_N)^{-1} y \quad \text{where } G_{ij} = \varphi(x_i)^\top \varphi(x_j) \text{ is the gram matrix}$$

**Kernel Ridge Regression** Essentially, the reformulation computes the similarity score between  $x$  and  $x'$ , which can be replaced by a kernel function  $K(x_i, x_j)$ , allowing computation in high-dimensional feature spaces without explicit

feature transformations. The solution to kernel ridge regression is:  $f(x) = \sum_{i=1}^N \alpha_i K(x_i, x)$  where  $\alpha_i$  are coefficients determined based on the training data and the kernel.

**Kernel Construction** To build a kernel, we can write an expression in the form of a dot product of the same function with variable  $x$  and  $x'$ , but this method does not always work.

**Mercer's Theorem and SPD Kernels** Suppose  $k$  is a SPD kernel. Then, there exists a feature space  $\mathcal{H}$  and a feature map  $\varphi: \mathbb{R}^d \rightarrow \mathcal{H}$  such that  $k(x, x') = \varphi(x)^\top \varphi(x')$

**SPD kernels properties:**

- **Symmetry:**  $K(x, x') = K(x', x)$
- **Positive Semi-definiteness:** For any  $n$  and  $\{x_1, \dots, x_n\}$ , the Gram matrix  $G_{ij} = k(x_i, x_j)$  is positive semi-definite. (Recall: a matrix  $G$  is positive semi-definite if  $c^T G c \geq 0$  for any vector  $c$ )

### Examples of SPD Kernels

- **Linear Kernel:**  $K(x, x') = x^\top x'$
- **Polynomial Kernel:**  $K(x, x') = (1 + x^\top x')^d$

$$\bullet \text{ Gaussian RBF Kernel: } K(x, x') = \exp\left(-\frac{\|x - x'\|^2}{2\sigma^2}\right)$$

**Constructing kernels** Given valid kernels  $k_1(x, x')$  and  $k_2(x, x')$ , the following new kernels will also be valid:  $k(x, x') = c k_1(x, x')$   $k(x, x') = f(x) k_1(x, x') f(x')$

$$k(x, x') = q(k_1(x, x')) \quad k(x, x') = \exp(k_1(x, x'))$$

$$k(x, x') = k_1(x, x') + k_2(x, x') \quad k(x, x') = k_1(x, x') k_2(x, x')$$

$$k(x, x') = k_3(\varphi(x), \varphi(x')) \quad k(x, x') = x^\top A x'$$

$$k(x, x') = k_a(x_a, x'_a) + k_b(x_b, x'_b) \quad k(x, x') = k_a(x_a, x'_a) k_b(x_b, x'_b)$$

### Lecture 3

#### Support Vector Machines (SVM)

Support Vector Machines (SVMs) are designed for binary classification. They find a hyperplane that separates two classes with the maximum margin, defined as the minimum distance between the separating hyperplane and the data points.

**Max Margin Formulation**  $\max_{w, b} \frac{1}{\|w\|} \min_{i=1, \dots, N} |w^T x_i + b|$  subject to  $y_i(w^T x_i + b) > 0 \quad \forall i$

**Optimization Problem**  $\min_{w, b} \frac{1}{2} \|w\|^2$  subject to:  $y_i(w^T x_i + b) \geq 1 \quad \forall i$ . Introducing Lagrange multipliers  $\alpha_i \geq 0$ , the Lagrangian is:  $\mathcal{L}(w, b, \alpha) = \frac{1}{2} \|w\|^2 - \sum_{i=1}^N \alpha_i [y_i(w^T x_i + b) - 1]$

**Karush-Kuhn-Tucker (KKT) Conditions:** Define the Lagrangian  $\mathcal{L}(z, \mu) = F(z) + \mu^T G(z)$ . Then, under technical conditions, for each locally optimal  $\hat{z}$ , there exists Lagrange multipliers  $\hat{\mu} \in \mathbb{R}^m$  such that:

**stationarity**  $\nabla_z \mathcal{L}(\hat{z}, \hat{\mu}) = 0$

**Primal Feasibility**  $G(\hat{z}) \leq 0$

**Dual Feasibility**  $\hat{\mu} \geq 0$

**Complementary Slackness**  $\hat{\mu}_j G_j(\hat{z}_j) = 0$

**Dual Problem**  $\max_{\mu \geq 0} \hat{F}(\mu)$  where  $\hat{F}(\mu) = \min_z \mathcal{L}(z, \mu) \mathcal{L}(z, \mu) = F(z) + \mu^T G(z)$

**KKT conditions for SVM**

1. **From Stationarity**  $\hat{w} = \sum_{i=1}^N \hat{\mu}_i y_i x_i, \quad 0 = \sum_{i=1}^N \hat{\mu}_i y_i$
2. **From Dual Feasibility**  $\hat{\mu}_i \geq 0$  for  $i = 1, \dots, N$

3. **From Complementary Slackness**  $\hat{\mu}_i = 0$  or  $y_i(\hat{w}^T x_i + \hat{b}) = 1$

4. **The multipliers**  $\hat{\mu}$  can be found by the dual problem  $\max_{\mu \in \mathbb{R}^N} \sum_{i=1}^N \mu_i - \frac{1}{2} \sum_{i,j} \mu_i \mu_j y_i y_j (x_i^T x_j)$

#### Dual formulation of SVM

$$\hat{\mu} = \arg \max_{\mu \in \mathbb{R}^N} \sum_{i=1}^N \mu_i - \frac{1}{2} \sum_{i,j} \mu_i \mu_j y_i y_j (x_i^T x_j)$$

Subject to:  $\hat{\mu} \geq 0$  and  $\sum_{i=1}^N \hat{\mu}_i y_i = 0$

Decision function:  $\hat{f}(x) = \text{sgn}\left(\sum_{i=1}^N \hat{\mu}_i y_i x_i^T x + \hat{b}\right)$

Complementary slackness:  $\hat{\mu}_i = 0$  or  $1 = y_i(\hat{w}^T x_i + \hat{b})$

#### Kernel Support Vector Machines

$$\hat{\mu} = \arg \max_{\mu \in \mathbb{R}^N} \sum_{i=1}^N \mu_i - \frac{1}{2} \sum_{i,j=1}^N \mu_i \mu_j y_i y_j k(x_i, x_j)$$

Subject to:  $\hat{\mu} \geq 0$  and  $\sum_{i=1}^N \hat{\mu}_i y_i = 0$

Decision function:  $\hat{f}(x) = \text{sgn}\left(\sum_{i=1}^N \hat{\mu}_i y_i k(x_i, x) + \hat{b}\right)$

**Only** support vectors satisfying  $1 = y_i(\hat{w}^T \varphi_i(x) + \hat{b})$  matter for predictions. This is a **sparse kernel method**.

## Lecture 4

**Decision Trees:** Stratify the input space into distinct, non-overlapping regions and assign a chosen, **constant** prediction to each region

**Universal approximation theorem** In general, any (sufficiently regular) oracle function  $f^* : \mathcal{X} \rightarrow \mathbb{R}$  can be approximated with a decision tree with an arbitrary small precision  $\epsilon > 0$ . i.e. there exists a decision tree  $T$  (with some finite depth) such that  $|T(x) - f^*(x)| < \epsilon$ . Want  $\epsilon$  to be small  $\Rightarrow$  choose larger  $L$

**Classification and Regression Trees** Suppose that the input space is  $\mathcal{X}$ . A **partition** of  $\mathcal{X}$  is a collection of subsets  $\mathcal{R}_1, \mathcal{R}_2, \dots, \mathcal{R}_J$  such that  $\mathcal{R}_i \cap \mathcal{R}_j = \emptyset$  for  $i \neq j$  and  $\bigcup_{j=1}^J \mathcal{R}_j = \mathcal{X}$

The general decision tree hypothesis space is:  $\mathcal{H} = \left\{ f : f(x) = \sum_{j=1}^J a_j \mathbb{1}_{x \in \mathcal{R}_j}, \{\mathcal{R}_j\} \text{ is a partition of } \mathcal{X}, a_j \in \mathbb{Y} \right\}$

where  $\mathbb{1}_{x \in \mathcal{R}_j} = \begin{cases} 1, & \text{if } x \in \mathcal{R}_j \\ 0, & \text{otherwise} \end{cases}$

### Learning Decision Trees

A decision tree model  $f(x) = \sum_{j=1}^J a_j \mathbb{1}_{x \in \mathcal{R}_j}$  depends on both  $a_j$  and  $\mathcal{R}_j$ . Given  $\{\mathcal{R}_j\}$ ,  $\{a_j\}$  are easy to fix:

- **Regression:** we take the **average** label values  $a_j = y_j = \frac{\sum_i y_i \mathbb{1}_{x \in \mathcal{R}_j}}{\sum_i \mathbb{1}_{x \in \mathcal{R}_j}}$
- **Classification:** we take the **modal** label values  $a_j = \text{mode}\{y_i : x_i \in \mathcal{R}_j\}$

Suppose we are dealing with regression, then we can fix  $a_j$  as before and solve the following empirical risk minimization:

$$\min_{\mathcal{R}_j} \frac{1}{2} \sum_{i=1}^N \left( \sum_{j=1}^J a_j \mathbb{1}_{x \in \mathcal{R}_j} - y_i \right)^2. \text{ However this is very hard to solve (NP-hard).}$$

**Recursive Binary Splitting Greedy** algorithm which essentially repeats the following two steps and grows the tree by adding two leaf nodes at a time: Pick a dimension of the input space, then find the **best** value  $\theta$  to split this input dimension into two parts and assign new constant values to these new regions.

**Loss function for Decision Trees Classification** Entropy:  $-\sum_{k=1}^K \sum_{j=1}^J p_{jk} \log p_{jk}$  Gini Impurity:  $\sum_{k=1}^K \sum_{j=1}^J p_{jk}(1 - p_{jk})$  where  $p_{jk}$  is the proportion of samples in  $\mathcal{R}_j$  belonging to class  $k$ .

**Advantages of Decision Trees:** 1. Can readily visualize and understand predictions; 2. Implicit feature selection via analyzing contribution of splits to reduction of error/impurity; 3. Robust to data types, supervised learning tasks, and nonlinear relationships.

**Disadvantages of Decision Trees:** 1. Greedy algorithms may find sub-optimal solutions; 2. Sensitive to data variation and balancing; 3. Prone to overfitting.

**Model Ensembling Bagging** reduces the variance by aggregating predictions from multiple models trained on different random subsamples of the data:

- **Regression**  $\hat{f}(x) = \frac{1}{m} \sum_{j=1}^m f_j(x)$
- **Classification**  $\hat{f}(x) = \text{Mode}\{f_j(x) : j = 1, \dots, m\}$

**Boosting** works by training weak learners sequentially, where each learner focuses on correcting the mistakes of the previous ones. The combined model is a weighted sum of the weak learners:  $f(x) = \sum_{t=1}^T \alpha_t f_t(x)$  where  $\alpha_t$  are coefficients based on each learner's performance. Boosting helps reduce bias.

**Key Ideas of AdaBoost** 1. Initialize with uniform weight across all training samples

2. Train a classifier/regressor  $f_1$
3. Identify the samples that  $f_1$  got wrong (classification) or has large errors (regression)
4. Weight these samples more heavily and train  $f_2$  on this reweighted dataset
5. Repeat steps 3-5

### AdaBoost Implementation

- **Data:**  $\mathcal{D} = \{(x_i, y_i)\}_{i=1}^N$
- Initialize  $w_i^{(1)} = \frac{1}{N}$  for all  $i = 1, \dots, N$ ;
- For  $j = 1, \dots, m$  do

1. Obtain  $f_j$  from:  $f_j = \arg \min_{f \in \mathcal{H}} \sum_{i=1}^N w_i^{(j)} \mathbb{1}_{y_i \neq f(x_i)}$
2. Compute combination coefficients:  $\delta_j = \frac{\sum_{i=1}^N w_i^{(j)} \mathbb{1}_{y_i \neq f_j(x_i)}}{\sum_{i=1}^N w_i^{(j)}} \quad \alpha_j = \log \left( \frac{1 - \delta_j}{\delta_j} \right)$
3. Update weights:  $w_i^{(j+1)} = w_i^{(j)} \exp \left( \alpha_j \mathbb{1}_{y_i \neq f_j(x_i)} \right)$

- Return:  $\bar{f}(x) = \text{Sign} \left( \sum_{j=1}^m \alpha_j f_j(x) \right)$

**Cross-Validation** Cross-validation is essential for tuning model hyperparameters. In  $k$ -fold cross-validation, the data is split into  $k$  subsets. The model is trained on  $k - 1$  subsets and validated on the remaining one. This process is repeated  $k$ -times, and the average performance is evaluated. Cross-validation ensures model generalizability.

## Neural Network Architecture Example

The architecture of the network consists of:  $x = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}$ ,  $x_1 \in \mathbb{R}$ ,  $x_2 \in \mathbb{R}$  **Input layer:**  $x_0 = x = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} \in \mathbb{R}^2$

**1st hidden layer with 2 neurons. 2nd hidden layer with 3 neurons. Output:** Scalar.

**1st Hidden Layer**  $x_1 = W_0 x_0 + b_0$   $W_0 \in \mathbb{R}^{2 \times 2}$ ,  $b_0 \in \mathbb{R}^{2 \times 1}$   $x_1 = \begin{bmatrix} x_{11} \\ x_{12} \end{bmatrix}$ ,  $x_1 \in \mathbb{R}^{2 \times 1}$   $x_1 = \text{ReLU}(x_1)$

**2nd Hidden Layer**  $x_2 = W_1 x_1 + b_1$   $W_1 \in \mathbb{R}^{3 \times 2}$ ,  $b_1 \in \mathbb{R}^{3 \times 1}$   $x_2 = \text{ReLU}(x_2)$ ,  $x_2 \in \mathbb{R}^{3 \times 1}$

### Back-propagation Example Computation

The model is defined as:  $y(x) = v \delta(w_1 \delta(w_0 x))$ ,  $x, w_0, w_1, v \in \mathbb{R}$ , where  $\delta$  is the identity function, i.e.,  $\delta(z) = z$ .

The loss function is:  $L = (y(x) - y)^2$  We want to compute the gradients:  $\frac{\partial L}{\partial v}$ ,  $\frac{\partial L}{\partial w_1}$ ,  $\frac{\partial L}{\partial w_0}$

**Forward pass:**  $x_1 = w_0 x$   $x_2 = v w_1 x_1$   $L = (x_2 - y)^2$

**Backward pass:**

$$p_2 = \frac{\partial L}{\partial x_2} = 2(x_2 - y) \quad \frac{\partial L}{\partial v} = \frac{\partial L}{\partial x_2} \cdot \frac{\partial x_2}{\partial v} = 2(x_2 - y) w_1 x_1 \quad \frac{\partial L}{\partial w_1} = \frac{\partial L}{\partial x_2} \cdot \frac{\partial x_2}{\partial w_1} = 2(x_2 - y) v x_1$$

$$p_1 = \frac{\partial L}{\partial x_1} = \frac{\partial L}{\partial x_2} \cdot \frac{\partial x_2}{\partial x_1} = 2(x_2 - y) v w_1 \quad \frac{\partial L}{\partial w_0} = \frac{\partial L}{\partial x_1} \cdot \frac{\partial x_1}{\partial w_0} = 2(x_2 - y) v w_1 x$$

## Lecture 5

**Bias** The difference between expected prediction and the true value.  $\text{bias}^2 = \int_{\mathbb{X}} \{\mathbb{E}_D[f(x)] - y\}^2 p(x) dx$

**Variance** The difference between what you expect to learn, i.e.,  $\bar{f}$ , and what you learn from a particular dataset.  $\text{variance} = \int_{\mathbb{X}} \mathbb{E}_D \left[ (f(x) - \bar{f}(x))^2 \right] p(x) dx$

**Bias-Variance Tradeoff** The bias-variance tradeoff is critical for understanding model performance:

$\text{MSE} = \text{Bias}^2 + \text{Variance} + \text{Irreducible Error}$

**Underfitting** High bias with low variance; **Overfitting** low bias with high variance.

**Neural Networks** The neural network (NN) hypothesis space is quite like linear basis models:  $f(x) = \sum_{j=1}^M v_j \phi_j(x)$ , where  $\phi_j(x) = \sigma(w_j^T x + b_j)$ ;  $w_j \in \mathbb{R}^d$ : the trainable weights of the hidden layer;  $b_j \in \mathbb{R}$ : the trainable biases of the hidden layer;  $v_j \in \mathbb{R}$ : the trainable weights of the output layer;  $\sigma : \mathbb{R} \rightarrow \mathbb{R}$  is the activation function.

### Activation Functions

**Sigmoid:**  $\sigma(z) = \frac{1}{1+e^{-z}}, [0, 1]$  **Tanh:**  $\sigma(z) = \tanh(z), [-1, 1]$

**Rectified Linear Unit (ReLU):**  $\sigma(z) = \max(0, z), [0, \infty]$  **Leaky-ReLU:**  $\sigma(z) = \begin{cases} z & \text{if } z \geq 0 \\ \delta z & \text{if } z < 0 \end{cases}, [-\infty, \infty]$

**Universal Approximation Theorem** Any continuous function  $f^*$  on a compact domain can be approximated by neural networks to arbitrary precision, provided there are enough neurons ( $M$  large enough).

**Curse of Dimensionality:** Exponential increase in number of neurons required as the dimensionality increases. Baron(1993) showed that for any continuous function  $f^* : [0, 1]^d \rightarrow \mathbb{R}$  (with other conditions), there exists a width- $M$  neural network  $f_M$  such that  $\|f^* - f_M\|^2 \leq \mathcal{O}(M^{-1})$

### The Significance of Data-Dependent Feature Maps

Functions behave just like vectors: Each  $\phi_j$  is like a coordinate axis. They play the role of  $e_j$ , but there are an infinite number of them. In this case the oracle function  $f^*$  plays the role of  $u$ . Writing  $f^*(x) = \sum_{j=1}^M v_j \phi_j(x)$  is like **expanding a vector into its components**, but we can't have all components since  $M$  is finite. If we get to choose which components to have in the sum **after** seeing some information on  $f^*$ , we can usually do much better (reduced error by choosing  $3e_1 + e_2$  and  $e_3$  is better than choosing raw bases  $e_1, e_2, e_3$ ).

**Gradient Descent**  $w^{(t+1)} = w^{(t)} - \eta \frac{\partial L}{\partial w}$  When  $\eta$  is too small, updates are **slow**; When  $\eta$  is too large, the updates may become **unstable**.

**Stochastic Gradient Descent (SGD):**  $\theta_{k+1} = \theta_k - \eta \frac{1}{B} \sum_{i \in I_B} \nabla \Phi_i(\theta_k)$ , this diminishes the probability of sticking at local minima as the main drawback for GD.

**Convex function** is a function satisfying:  $\Phi(\lambda\theta + (1-\lambda)\theta') \leq \lambda\Phi(\theta) + (1-\lambda)\Phi(\theta')$  for all  $\theta, \theta' \in \mathbb{R}^p$  and  $\lambda \in [0, 1]$

### Back-propagation Algorithm

- Initialize  $x_0 = x \in \mathbb{R}^d$ .
- For  $t = 0, 1, \dots, T$ :  $x_{t+1} = g_t(x_t, W_t) = \sigma(W_t x_t)$
- Set  $p_{T+1} = \nabla_x L(x_{T+1}, y)$ .
- For  $t = T, T-1, \dots, 1$ :
  - $\nabla W_t \Phi = p_{t+1}^T \nabla W g_t(x_t, W_t)$
  - $p_t = [\nabla_x g_t(x_t, W_t)]^T p_{t+1}$
- Return  $\{\nabla W_t \Phi : t = 0, \dots, T\}$ .

## Unsupervised Learning

### Lecture 6

**Principal Component Analysis** Two formulations:

- Find the direction that captures the most variance, where the order in magnitude of the normalized sample covariance matrix  $S$  corresponds to the order of eigenvalues.  $\lambda$
- Find the direction that minimizes projection error by partitioning the set of eigenvectors and choosing the exclusion set ( $m+1$  to  $d$ ) to be the eigenvector subset corresponding to the smallest ( $d - m - 1$ ) eigenvalues.

### PCA Algorithm Implementation

Center data  $\rightarrow$  compute sample covariance matrix  $\rightarrow$  compute eigenvectors of largest eigenvalues

- **Data:**  $\mathcal{D} = \{x_i\}_{i=1}^N$ ,  $x_i \in \mathbb{R}^d$  for all  $i$ .
- **Hyper-parameters:**  $m$  (reduction dimension).
- Compute sample covariance matrix  $S = \frac{1}{N} \sum_{i=1}^N x_i x_i^T$ .
- Compute the first  $m$  eigenvectors  $\{u_1, \dots, u_m\}$  and eigenvalues  $\{\lambda_1, \dots, \lambda_m\}$ .
- Form  $d \times m$  matrix  $U_m$  whose  $j^{\text{th}}$  column is  $u_j$ .
- Compute  $Z_m = XU_m$ .
- Return Principal component scores  $Z_m$ , Eigenvalues, and eigenvectors  $\lambda_j, u_j$  for  $j = 1, \dots, m$ .

### PCA in Feature Space

- **Data:**  $\mathcal{D} = \{x_i\}_{i=1}^N$ ,  $x_i \in \mathbb{R}^d$  for all  $i$
- **Hyper-parameters:**  $m$  (reduction dimension),  $\phi$  (feature maps);
- Compute design matrix  $\Phi_{ij} = \phi_j(x_i)$ ;
- Center design matrix  $\Phi_{ij} \leftarrow \Phi_{ij} - \frac{1}{N} \sum_{i=1}^N \Phi_{ij}$ ;
- Compute sample covariance matrix  $S_\phi = \frac{1}{N} \Phi^T \Phi$ ;
- Compute the first  $m$  eigenvectors  $\{u_1, \dots, u_m\}$  and eigenvalues  $\{\lambda_1, \dots, \lambda_m\}$  of  $S_\phi$ ;
- Form  $d \times m$  matrix  $U_m$  whose  $j^{\text{th}}$  column is  $u_j$ ;
- Compute  $Z_m = \Phi U_m$ ;
- **Return:** Principal component score  $Z_m$ , Eigenvalues and eigenvectors  $\lambda_j, u_j$  for  $j = 1, \dots, m$ .

**PCA whitening transform** Principal component scores are given by  $Z = XU$  where  $X$  is the original features and  $U$  is the matrix of eigenvectors. The transformation  $X' = XU\Lambda^{-\frac{1}{2}}$ , where  $\Lambda$  is the matrix of eigenvalues makes  $\text{cov}(X') = I$ .

### Useful derivatives

$$\frac{\partial}{\partial x} (x^T A x) = 2Ax \quad \frac{\partial}{\partial x} (A^T x) = A$$