

# DSA5105 Principles of Machine Learning

AY2024/25 Sem1 By Zhao Peiduo

## Supervised Learning

### Lecture 1

**Empirical Risk Minimization (ERM)** The learning process aims to find a function  $f \in \mathcal{H}$  that minimizes the empirical risk,  $R_{\text{emp}}(f) = \frac{1}{N} \sum_{i=1}^N L(y_i, f(x_i))$ , where  $y_i = f^*(x_i)$ .

**Loss Function** To quantify how well a function  $f$  fits the data, we use a loss function  $L(y, \hat{y})$ , where  $y$  is the true output, and  $\hat{y} = f(x)$  is the predicted output. Common loss functions include the mean squared error (MSE) for regression tasks:  $L(y, \hat{y}) = \frac{1}{2}(y - \hat{y})^2$ , cross-entropy loss for classification task:  $L(y, p) = -\sum_i y_i \log(p_i)$ , and huber

loss for robust regression:  $L_\delta(y, \hat{y}) = \begin{cases} \frac{1}{2}(y - \hat{y})^2 & \text{if } |y - \hat{y}| \leq \delta \\ \delta|y - \hat{y}| - \frac{1}{2}\delta^2 & \text{otherwise} \end{cases}$

**Softmax Function** For a multi-class classification problem with  $K$  classes, the softmax function is defined as:  $\text{softmax}(z_i) = \frac{\exp(z_i)}{\sum_{j=1}^K \exp(z_j)}$

**General Ordinary Least Squares Formula** Consider  $x \in \mathbb{R}^d$  and the new hypothesis space  $\mathcal{H}_M = \{f : f(x) = \sum_{j=0}^{M-1} w_j \varphi_j(x)\}$ . Each  $\varphi_j : \mathbb{R}^d \rightarrow \mathbb{R}$  is called a **basis function** or **feature map**.

We can rewrite the ERM  $\min_{w \in \mathbb{R}^M} \frac{1}{2N} \sum_{i=1}^N \left( \sum_{j=0}^{M-1} w_j \varphi_j(x_i) - y_i \right)^2$  into  $\min_{w \in \mathbb{R}^M} \frac{1}{2N} \|\Phi w - y\|^2$ . Solving by setting  $\nabla R_{\text{emp}}(\hat{w}) = 0$ , we have  $\hat{w} = (\Phi^T \Phi)^{-1} \Phi^T y$ , given invertible  $\Phi^T \Phi$ .

For cases where  $\Phi^T \Phi$  is not invertible, the formula using the Moore-Penrose pseudoinverse is:  $\hat{w}(u) = \Phi^\dagger y + (I - \Phi^\dagger \Phi)u$ ,  $u \in \mathbb{R}^M$ .

**Regularization:** To prevent overfitting, regularization techniques add a penalty to the loss function:  $\min_{w \in \mathbb{R}^M} \frac{1}{2N} \|\Phi w - y\|^2 + \lambda C(w)$ ,  $\frac{1}{2N} \|\Phi w - y\|^2 + \lambda C(w)$ . Minimizing the ERM we get  $\hat{w} = (\Phi^T \Phi + \lambda I_M)^{-1} \Phi^T y$  which is always invertible for positive  $\lambda$ .

Common regularization terms:  $L_2$  (Ridge) regularization:  $\lambda \sum_{j=1}^p w_j^2$ ,  $L_1$  (Lasso) regularization:  $\lambda \sum_{j=1}^p |w_j|$

### Lecture 2

**Reformulation of Ridge Regression** We rewrite the regularized least squares solution in another way:

$$\hat{w} = (\Phi^T \Phi + \lambda I_M)^{-1} \Phi^T y = \Phi^T (\Phi \Phi^T + \lambda I_N)^{-1} y$$

#### Reformulation of Ridge Regression

$f(x) = \sum_{i=1}^N \alpha_i \varphi(x_i)^T \varphi(x)$ ,  $\alpha = (G + \lambda I_N)^{-1} y$  where  $G_{ij} = \varphi(x_i)^T \varphi(x_j)$  is the gram matrix

**Kernel Ridge Regression** Essentially, the reformulation computes the similarity score between  $x$  and  $x'$ , which can be replaced by a kernel function  $K(x_i, x_j)$ , allowing computation in high-dimensional feature spaces without explicit

feature transformations. The solution to kernel ridge regression is:  $f(x) = \sum_{i=1}^N \alpha_i K(x_i, x)$  where  $\alpha_i$  are coefficients determined based on the training data and the kernel.

**Mercer's Theorem and SPD Kernels** Suppose  $k$  is a SPD kernel. Then, there exists a feature space  $\mathcal{H}$  and a feature map  $\varphi : \mathbb{R}^d \rightarrow \mathcal{H}$  such that  $k(x, x') = \varphi(x)^T \varphi(x')$

**SPD kernels properties:**  $K(x, x') = K(x', x)$  (**Symmetry**) For any  $n$  and  $\{x_1, \dots, x_n\}$ , the Gram matrix  $G_{ij} = k(x_i, x_j)$  is positive semi-definite. (Recall: a matrix  $G$  is positive semi-definite if  $c^T G c \geq 0$  for any vector  $c$ ) (**Positive Semi-definiteness**)

**Examples of SPD Kernels:** **Linear Kernel:**  $K(x, x') = x^T x'$  **Polynomial Kernel:**  $K(x, x') = (1 + x^T x')^d$

**Gaussian RBF Kernel:**  $K(x, x') = \exp\left(-\frac{\|x - x'\|^2}{2\sigma^2}\right)$

**Constructing kernels** Given valid kernels  $k_1(x, x')$  and  $k_2(x, x')$ , the following new kernels will also be valid:  $k(x, x') = ck_1(x, x')$   $k(x, x') = f(x)k_1(x, x')f(x')$   $k(x, x') = q(k_1(x, x'))$   $k(x, x') = \exp(k_1(x, x'))$

$$k(x, x') = k_1(x, x') + k_2(x, x') \quad k(x, x') = k_1(x, x')k_2(x, x') \quad k(x, x') = k_3(\varphi(x), \varphi(x'))$$

$$k(x, x') = x^T A x' \quad k(x, x') = k_a(x_a, x'_a) + k_b(x_b, x'_b) \quad k(x, x') = k_a(x_a, x'_a)k_b(x_b, x'_b)$$

### Lecture 3

**SVM Max Margin Formulation**  $\max_{w, b} \frac{1}{\|w\|} \min_{i=1, \dots, N} |w^T x_i + b|$  subject to  $y_i(w^T x_i + b) > 0 \quad \forall i$

**Optimization Problem**  $\min_{w, b} \frac{1}{2} \|w\|^2$  subject to:  $y_i(w^T x_i + b) \geq 1 \quad \forall i$ . Introducing Lagrange multipliers  $\alpha_i \geq 0$ , the Lagrangian is:  $\mathcal{L}(w, b, \alpha) = \frac{1}{2} \|w\|^2 - \sum_{i=1}^N \alpha_i [y_i(w^T x_i + b) - 1]$

**Karush-Kuhn-Tucker (KKT) Conditions:** Define the Lagrangian  $\mathcal{L}(z, \mu) = F(z) + \mu^T G(z)$ . Then, under technical conditions, for each locally optimal  $\hat{z}$ , there exists Lagrange multipliers  $\hat{\mu} \in \mathbb{R}^m$  such that:

**stationarity**  $\nabla_z \mathcal{L}(\hat{z}, \hat{\mu}) = 0$  **Primal Feasibility**  $G(\hat{z}) \leq 0$

**Dual Feasibility**  $\hat{\mu} \geq 0$  **Complementary Slackness**  $\hat{\mu}_j G_j(\hat{z}_j) = 0$

**Dual Problem**  $\max_{\mu} \mu \geq 0 \quad \hat{F}(\mu)$  where  $\hat{F}(\mu) = \min_z \mathcal{L}(z, \mu)$ ,  $\mathcal{L}(z, \mu) = F(z) + \mu^T G(z)$

**KKT conditions for SVM** [noitemsep, topsep=0pt]

1. From Stationarity  $\hat{w} = \sum_{i=1}^N \hat{\mu}_i y_i x_i$ ,  $0 = \sum_{i=1}^N \hat{\mu}_i y_i$

2. From Dual Feasibility  $\hat{\mu}_i \geq 0$  for  $i = 1, \dots, N$

3. From Complementary Slackness  $\hat{\mu}_i = 0$  or  $y_i(\hat{w}^T x_i + \hat{b}) = 1$

4. The multipliers  $\hat{\mu}$  can be found by the dual problem  $\max_{\mu \in \mathbb{R}^N} \sum_{i=1}^N \mu_i - \frac{1}{2} \sum_{i,j} \mu_i \mu_j y_i y_j (x_i^T x_j)$

**Dual formulation of SVM**

$\hat{\mu} = \arg \max_{\mu \in \mathbb{R}^N} \sum_{i=1}^N \mu_i - \frac{1}{2} \sum_{i,j} \mu_i \mu_j y_i y_j (x_i^T x_j)$  Subject to:  $\hat{\mu} \geq 0$  and  $\sum_{i=1}^N \hat{\mu}_i y_i = 0$

Decision function:  $\hat{f}(x) = \text{sgn}\left(\sum_{i=1}^N \hat{\mu}_i y_i x_i^T x + \hat{b}\right)$  Complementary slackness:  $\hat{\mu}_i = 0$  or  $1 = y_i(\hat{w}^T x_i + \hat{b})$

**Kernel SVMs**

$\hat{\mu} = \arg \max_{\mu \in \mathbb{R}^N} \sum_{i=1}^N \mu_i - \frac{1}{2} \sum_{i,j=1}^N \mu_i \mu_j y_i y_j k(x_i, x_j)$  Subject to:  $\hat{\mu} \geq 0$  and  $\sum_{i=1}^N \hat{\mu}_i y_i = 0$

Decision function:  $\hat{f}(x) = \text{sgn}\left(\sum_{i=1}^N \hat{\mu}_i y_i k(x_i, x) + \hat{b}\right)$

Only support vectors satisfying  $1 = y_i(\hat{w}^T \varphi_i(x) + \hat{b})$  matter for predictions. This is a **sparse kernel method**.

### Lecture 4

**Classification and Regression Trees** Suppose that the input space is  $\mathcal{X}$ . A **partition** of  $\mathcal{X}$  is a collection of subsets  $\mathcal{R}_1, \mathcal{R}_2, \dots, \mathcal{R}_J$  such that  $\mathcal{R}_i \cap \mathcal{R}_j = \emptyset$  for  $i \neq j$  and  $\bigcup_{j=1}^J \mathcal{R}_j = \mathcal{X}$

The general decision tree hypothesis space is:  $\mathcal{H} = \left\{f : f(x) = \sum_{j=1}^J a_j \mathbb{I}_{x \in \mathcal{R}_j}, \{\mathcal{R}_j\} \text{ is a partition of } \mathcal{X}, a_j \in \mathcal{Y}\right\}$

where  $\mathbb{I}_{x \in \mathcal{R}_j} = \begin{cases} 1, & \text{if } x \in \mathcal{R}_j \\ 0, & \text{otherwise} \end{cases}$

A decision tree model  $f(x) = \sum_{j=1}^J a_j \mathbb{I}_{x \in \mathcal{R}_j}$  depends on both  $a_j$  and  $\mathcal{R}_j$ . For regression we take the **average** label

values  $a_j = y_j = \frac{\sum_i y_i \mathbb{I}_{x \in \mathcal{R}_j}}{\sum_i \mathbb{I}_{x \in \mathcal{R}_j}}$ ; For classification we take the **mode** label values  $a_j = \text{mode}\{y_i : x_i \in \mathcal{R}_j\}$

**Loss function for Decision Trees Classification** Entropy:  $-\sum_{k=1}^K \sum_{j=1}^J p_{jk} \log p_{jk}$

Gini Impurity:  $\sum_{k=1}^K \sum_{j=1}^J p_{jk}(1 - p_{jk})$  where  $p_{jk}$  is the proportion of samples in  $\mathcal{R}_j$  belonging to class  $k$ .

**Model Ensembling Bagging** reduces the variance by aggregating predictions from multiple models trained on different random subsamples of the data: **Regression**  $\hat{f}(x) = \frac{1}{m} \sum_{j=1}^m f_j(x)$  **Classification**  $\hat{f}(x) = \text{Mode}\{f_j(x) : j = 1, \dots, m\}$

**Boosting** works by training weak learners sequentially, where each learner focuses on correcting the mistakes of the previous ones. The combined model is a weighted sum of the weak learners:  $f(x) = \sum_{t=1}^T \alpha_t f_t(x)$  where  $\alpha_t$  are coefficients based on each learner's performance. Boosting helps reduce bias.

**Key Ideas of AdaBoost** 1. Initialize with uniform weight across all training samples

2. Train a classifier/regressor  $f_1$

3. Identify the samples that  $f_1$  got wrong (classification) or has large errors (regression)

4. Weight these samples more heavily and train  $f_2$  on this reweighted dataset

5. Repeat steps 3-5

**AdaBoost Implementation**

- **Data:**  $\mathcal{D} = \{(x_i, y_i)\}_{i=1}^N$
- Initialize  $w_i^{(1)} = \frac{1}{N}$  for all  $i = 1, \dots, N$ ;
- For  $j = 1, \dots, m$  do

1. Obtain  $f_j$  from:  $f_j = \arg \min_{f \in \mathcal{H}} \sum_{i=1}^N w_i^{(j)} \mathbb{I}_{y_i \neq f(x_i)}$

2. Compute combination coefficients:  $\delta_j = \frac{\sum_{i=1}^N w_i^{(j)} \mathbb{I}_{y_i \neq f_j(x_i)}}{\sum_{i=1}^N w_i^{(j)}}$   $\alpha_j = \log\left(\frac{1 - \delta_j}{\delta_j}\right)$

3. Update weights:  $w_i^{(j+1)} = w_i^{(j)} \exp\left(\alpha_j \mathbb{I}_{y_i \neq f_j(x_i)}\right)$

- Return:  $\hat{f}(x) = \text{Sign}\left(\sum_{j=1}^m \alpha_j f_j(x)\right)$

**Cross-Validation** Cross-validation is essential for tuning model hyperparameters. In  $k$ -fold cross-validation, the data is split into  $k$  subsets. The model is trained on  $k - 1$  subsets and validated on the remaining one. This process is repeated  $k$ -times, and the average performance is evaluated. Cross-validation ensures model generalizability.

### Lecture 5

#### Activation Functions

**Sigmoid:**  $\sigma(z) = \frac{1}{1 + e^{-z}}$ ,  $[0, 1]$  **Tanh:**  $\sigma(z) = \tanh(z)$ ,  $[-1, 1]$  **ReLU**  $\sigma(z) = \max(0, z)$ ,  $[0, \text{inf}]$ , leaky ReLU:  $\delta \text{ if } z < 0$  instead of 0

**Gradient Descent**  $w^{(t+1)} = w^{(t)} - \eta \frac{\partial L}{\partial w}$  When  $\eta$  is too small, updates are **slow**; When  $\eta$  is too large, the updates may become **unstable**.

**Stochastic Gradient Descent (SGD):**  $\theta_{k+1} = \theta_k - \eta \frac{1}{B} \sum_{i \in I_B} \nabla \Phi_i(\theta_k)$ , this diminishes the probability of sticking at local minima as the main drawback for GD.

**Convex function** is a function satisfying:  $\Phi(\lambda\theta + (1 - \lambda)\theta') \leq \lambda\Phi(\theta) + (1 - \lambda)\Phi(\theta')$  for all  $\theta, \theta' \in \mathbb{R}^p$  and  $\lambda \in [0, 1]$  **Deep Neural Networks (DNNs)** Deep neural networks are an extension of shallow networks. The idea is to stack  $n$  hidden layers together and forward pass the  $x$  sequentially.

#### Back-propagation Algorithm

- Initialize  $x_0 = x \in \mathbb{R}^d$ .
- For  $t = 0, 1, \dots, T$ :  $x_{t+1} = g_t(x_t, W_t) = \sigma(W_t x_t)$
- Set  $p_{T+1} = \nabla_x L(x_{T+1}, y)$ .
- For  $t = T, T - 1, \dots, 1$ :

$$- \nabla_{W_t} \Phi = p_{t+1}^T \nabla_{W_t} g_t(x_t, W_t)$$

$$- p_t = [\nabla_x g_t(x_t, W_t)]^T p_{t+1}$$

- Return  $\{\nabla_{W_t} \Phi : t = 0, \dots, T\}$ .

#### Neural Network Architecture Example

**The architecture of the network consists of:**  $x = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}$ ,  $x_1 \in \mathbb{R}$ ,  $x_2 \in \mathbb{R}$  **Input layer:**  $x_0 = x = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} \in \mathbb{R}^2$

**1st hidden layer with 2 neurons. 2nd hidden layer with 3 neurons. Output:** Scalar.

**1st Hidden Layer**  $x_1 = W_0 x_0 + b_0$   $W_0 \in \mathbb{R}^{2 \times 2}$ ,  $b_0 \in \mathbb{R}^{2 \times 1}$   $x_1 = \begin{bmatrix} x_{11} \\ x_{12} \end{bmatrix}$ ,  $x_1 \in \mathbb{R}^{2 \times 1}$   $x_1 = \text{ReLU}(x_0)$

**2nd Hidden Layer**  $x_2 = W_1 x_1 + b_1$   $W_1 \in \mathbb{R}^{3 \times 2}$ ,  $b_1 \in \mathbb{R}^{3 \times 1}$   $x_2 = \text{ReLU}(x_1)$ ,  $x_2 \in \mathbb{R}^{3 \times 1}$

#### Back-propagation Example Computation

The model is defined as:  $y(x) = v\delta(w_1\delta(w_0x))$ ,  $x, w_0, w_1, v \in \mathbb{R}$ , where  $\delta$  is the identity function, i.e.,  $\delta(z) = z$ . The loss function is:  $L = (y(x) - y)^2$  We want to compute the gradients:  $\frac{\partial L}{\partial v}$ ,  $\frac{\partial L}{\partial w_1}$ ,  $\frac{\partial L}{\partial w_0}$

**Forward pass:**  $x_1 = w_0 x$   $x_2 = v w_1 x_1$   $L = (x_2 - y)^2$

**Backward pass:**

$$p_2 = \frac{\partial L}{\partial x_2} = 2(x_2 - y) \quad \frac{\partial L}{\partial v} = \frac{\partial L}{\partial x_2} \cdot \frac{\partial x_2}{\partial v} = 2(x_2 - y)w_1 x_1 \quad \frac{\partial L}{\partial w_1} = \frac{\partial L}{\partial x_2} \cdot \frac{\partial x_2}{\partial w_1} = 2(x_2 - y)v x_1$$

$$p_1 = \frac{\partial L}{\partial x_1} = \frac{\partial L}{\partial x_2} \cdot \frac{\partial x_2}{\partial x_1} = 2(x_2 - y)v w_1 \quad \frac{\partial L}{\partial w_0} = \frac{\partial L}{\partial x_1} \cdot \frac{\partial x_1}{\partial w_0} = 2(x_2 - y)v w_1 x$$

## Unsupervised Learning

### Lecture 6

#### PCA Algorithm Implementation

- **Data:**  $\mathcal{D} = \{x_i\}_{i=1}^N$ ,  $x_i \in \mathbb{R}^d$  for all  $i$ .
- **Hyper-parameters:**  $m$  (reduction dimension).
- Compute sample covariance matrix  $S = \frac{1}{N} \sum_{i=1}^N x_i x_i^T$ .
- Compute the first  $m$  eigenvectors  $\{u_1, \dots, u_m\}$  and eigenvalues  $\{\lambda_1, \dots, \lambda_m\}$ .
- Form  $d \times m$  matrix  $U_m$  whose  $j^{th}$  column is  $u_j$ .
- Compute  $Z_m = XU_m$ .
- Return Principal component scores  $Z_m$ , Eigenvalues, and eigenvectors  $\lambda_j, u_j$  for  $j = 1, \dots, m$ .

#### PCA in Feature Space

- **Data:**  $\mathcal{D} = \{x_i\}_{i=1}^N$ ,  $x_i \in \mathbb{R}^d$  for all  $i$
- **Hyper-parameters:**  $m$  (reduction dimension),  $\phi$  (feature maps);
- Compute design matrix  $\Phi_{ij} = \phi_j(x_i)$ ;
- Center design matrix  $\Phi_{ij} \leftarrow \Phi_{ij} - \frac{1}{N} \sum_{i=1}^N \Phi_{ij}$ ;
- Compute sample covariance matrix  $S_\phi = \frac{1}{N} \Phi \Phi^\top$ ;
- Compute the first  $m$  eigenvectors  $\{u_1, \dots, u_m\}$  and eigenvalues  $\{\lambda_1, \dots, \lambda_m\}$  of  $S_\phi$ ;
- Form  $d \times m$  matrix  $U_m$  whose  $j^{th}$  column is  $u_j$ ;
- Compute  $Z_m = \Phi U_m$ ;
- **Return:** Principal component score  $Z_m$ , Eigenvalues and eigenvectors  $\lambda_j, u_j$  for  $j = 1, \dots, m$ .

**PCA whitening transform** Principal component scores are given by  $Z = XU$  where  $X$  is the original features and  $U$  is the matrix of eigenvectors. The transformation  $X' = XU\Lambda^{-\frac{1}{2}}$ , where  $\Lambda$  is the matrix of eigenvalues makes  $\text{cov}(X') = I$ .

#### Useful derivatives

$$\frac{\partial}{\partial x} (x^T A x) = 2Ax \quad \frac{\partial}{\partial x} (A^T x) = A$$