

DSA5105 Principles of Machine Learning

AY2024/25 Sem1 By Zhao Peiduo

Supervised Learning

Lecture 1

Supervised Learning Supervised learning is the most common type of machine learning problem, where the goal is to make predictions and learn a function that maps an input to an output based on example input-output pairs.

Problem Setup Given a dataset $\mathcal{D} = \{(x_i, y_i)\}_{i=1}^N$, where x_i are the inputs and y_i are the corresponding outputs or labels, the goal is to learn the mapping $f: X \rightarrow Y$ such that $f(x_i) \approx y_i$ for all i .

Hypothesis Space The **oracle** f^* is unknown to us except through the dataset. The function f is chosen from a hypothesis space \mathcal{H} , which is a set of candidate functions. For example, in linear regression, the hypothesis space might be $\mathcal{H} = \{f: f(x) = w_0 + w_1 x \mid w_0, w_1 \in \mathbb{R}\}$.

Three Paradigms of Supervised Learning

- **Approximation:** Analyze the breadth and depth of the hypothesis space \mathcal{H} to determine if it contains, or closely approximates, the optimal function. (How large is our hypothesis space?)
- **Optimization:** Design and implement efficient algorithms to address the empirical risk minimization problem and find or closely approximate the best function within \mathcal{H} . (How can we find or get close to an approximation \hat{f} of f^* ?)
- **Generalization:** Evaluate whether the optimized model can effectively generalize to new, unseen data, focusing on the interplay between data size and model complexity. This is done through minimizing the population risk: $R_{\text{pop}}(f) = \mathbb{E}_{(x,y) \sim P}[L(y, f(x))]$. (Can the \hat{f} found generalized to unseen examples?)

Empirical Risk Minimization (ERM) The learning process aims to find a function $f \in \mathcal{H}$ that minimizes the empirical risk, $R_{\text{emp}}(f) = \frac{1}{N} \sum_{i=1}^N L(y_i, f(x_i))$, where $y_i = f^*(x_i)$.

Loss Function To quantify how well a function f fits the data, we use a loss function $L(y, \hat{y})$, where y is the true output, and $\hat{y} = f(x)$ is the predicted output. Common loss functions include the mean squared error (MSE) for regression tasks: $L(y, \hat{y}) = \frac{1}{2} (y - \hat{y})^2$

Ordinary Least Squares Formula The formula for Ordinary Least Squares in a simple linear regression context is given by: $\beta = (X^T X)^{-1} X^T y$. In 1D context, the empirical risk minimization problem can be defined as $\min_{w_0, w_1} \frac{1}{2N} \sum_{i=1}^N (w_0 + w_1 x_i - y_i)^2$. By setting the partial derivatives to zero, the ordinary least squares formula 1D is given by: $\hat{w}_0 = \bar{y} - \hat{w}_1 \bar{x}$ $\hat{w}_1 = \frac{\sum_i (x_i - \bar{x})(y_i - \bar{y})}{\sum_i (x_i - \bar{x})^2}$ $\bar{x} = \frac{1}{N} \sum_i x_i$ $\bar{y} = \frac{1}{N} \sum_i y_i$

Huber Loss The Huber loss, used for robust regression, is defined as: $L_\delta(y, \hat{y}) = \begin{cases} \frac{1}{2} (y - \hat{y})^2 & \text{if } |y - \hat{y}| \leq \delta \\ \delta |y - \hat{y}| - \frac{1}{2} \delta^2 & \text{otherwise} \end{cases}$

General Ordinary Least Squares Formula Consider $x \in \mathbb{R}^d$ and the new hypothesis space $\mathcal{H}_M = \{f: f(x) = \sum_{j=0}^{M-1} w_j \varphi_j(x)\}$ Each $\varphi_j: \mathbb{R}^d \rightarrow \mathbb{R}$ is called a **basis function** or **feature map**.

We can rewrite the ERM $\min_{w \in \mathbb{R}^M} \frac{1}{2N} \sum_{i=1}^N \left(\sum_{j=0}^{M-1} w_j \varphi_j(x_i) - y_i \right)^2$ into $\min_{w \in \mathbb{R}^M} \frac{1}{2N} \|\Phi w - y\|^2$. Solving by setting $\nabla R_{\text{emp}}(\hat{w}) = 0$, we have $\hat{w} = (\Phi^T \Phi)^{-1} \Phi^T y$, given invertible $\Phi^T \Phi$.

For cases where $\Phi^T \Phi$ is not invertible, the formula using the Moore-Penrose pseudoinverse is: $\hat{w}(u) = \Phi^\dagger y + (I - \Phi^\dagger \Phi) u$ $u \in \mathbb{R}^M$

Overfitting: Overfitting occurs when the hypothesis space \mathcal{H} is too large, allowing the model to fit the noise in the training data. This results in poor generalization to new data.

Regularization: To prevent overfitting, regularization techniques add a penalty to the loss function: $\min_{w \in \mathbb{R}^M} \frac{1}{2N} \|\Phi w - y\|^2 + \lambda C(w)$ $\frac{1}{2N} \|\Phi w - y\|^2 + \lambda C(w)$. Minimizing the ERM we get $\hat{w} = (\Phi^T \Phi + \lambda I_M)^{-1} \Phi^T y$ which is always invertible for positive λ .

Common regularization terms:

- **L₂ (Ridge) regularization:** $\lambda \sum_{j=1}^p w_j^2$
- **L₁ (Lasso) regularization:** $\lambda \sum_{j=1}^p |w_j|$

Softmax Function For a multi-class classification problem with K classes, the softmax function is defined as: $\text{softmax}(z_i) = \frac{\exp(z_i)}{\sum_{j=1}^K \exp(z_j)}$

Cross-Entropy Loss Commonly used in classification tasks, the cross-entropy loss is: $L(y, p) = - \sum_i y_i \log(p_i)$

Lecture 2

Another view of feature maps One can also view feature maps as implicitly defining some sort of **similarity measure**. Consider two vectors u and v . Then, $u^T v$ measures how similar they are. Feature maps define a **similarity** between two samples x, x' by computing the dot product in **feature space**.

Reformulation of Ridge Regression We rewrite the regularized least squares solution in another way:

$$\hat{w} = (\Phi^T \Phi + \lambda I_M)^{-1} \Phi^T y = \Phi^T (\Phi \Phi^T + \lambda I_N)^{-1} y$$

Proof:

$$(\Phi^T \Phi + \lambda I_M) \Phi^T = \Phi^T (\Phi \Phi^T + \lambda I_N)^{-1}$$

$$= \Phi^T (\Phi \Phi^T + \lambda I_N)^{-1} y$$

$$= \Phi^T \Phi^T y (\Phi^T + \lambda I_N)^{-1}$$

$$\lambda I_N^{-1} y = (\Phi^T \Phi + \lambda I_M)^{-1} \Phi^T y$$

Reformulation of Ridge Regression

$$\hat{f}(x) = \sum_{i=1}^N \alpha_i \varphi(x_i)^\top \varphi(x)$$

$$\alpha = (G + \lambda I_N)^{-1} y \quad \text{where } G_{ij} = \varphi(x_i)^\top \varphi(x_j) \text{ is the gram matrix}$$

Kernel Ridge Regression Essentially, the reformulation computes the similarity score between x and x' , which can be replaced by a kernel function $K(x_i, x_j)$, allowing computation in high-dimensional feature spaces without explicit

feature transformations. The solution to kernel ridge regression is: $f(x) = \sum_{i=1}^N \alpha_i K(x_i, x)$ where α_i are coefficients determined based on the training data and the kernel.

Kernel Construction To build a kernel, we can write an expression in the form of a dot product of the same function with variable x and x' , but this method does not always work.

Mercer's Theorem and SPD Kernels Suppose k is a SPD kernel. Then, there exists a feature space \mathcal{H} and a feature map $\varphi: \mathbb{R}^d \rightarrow \mathcal{H}$ such that $k(x, x') = \varphi(x)^\top \varphi(x')$

SPD kernels properties:

- **Symmetry:** $K(x, x') = K(x', x)$
- **Positive Semi-definiteness:** For any n and $\{x_1, \dots, x_n\}$, the Gram matrix $G_{ij} = k(x_i, x_j)$ is positive semi-definite. (Recall: a matrix G is positive semi-definite if $c^T G c \geq 0$ for any vector c)

Examples of SPD Kernels

- **Linear Kernel:** $K(x, x') = x^\top x'$
- **Polynomial Kernel:** $K(x, x') = (1 + x^\top x')^d$

$$\bullet \text{ Gaussian RBF Kernel: } K(x, x') = \exp\left(-\frac{\|x - x'\|^2}{2\sigma^2}\right)$$

Constructing kernels Given valid kernels $k_1(x, x')$ and $k_2(x, x')$, the following new kernels will also be valid: $k(x, x') = c k_1(x, x')$ $k(x, x') = f(x) k_1(x, x') f(x')$

$$k(x, x') = q(k_1(x, x')) \quad k(x, x') = \exp(k_1(x, x'))$$

$$k(x, x') = k_1(x, x') + k_2(x, x') \quad k(x, x') = k_1(x, x') k_2(x, x')$$

$$k(x, x') = k_3(\varphi(x), \varphi(x')) \quad k(x, x') = x^\top A x'$$

$$k(x, x') = k_a(x_a, x'_a) + k_b(x_b, x'_b) \quad k(x, x') = k_a(x_a, x'_a) k_b(x_b, x'_b)$$

Lecture 3

Support Vector Machines (SVM)

Support Vector Machines (SVMs) are designed for binary classification. They find a hyperplane that separates two classes with the maximum margin, defined as the minimum distance between the separating hyperplane and the data points.

Max Margin Formulation $\max_{w, b} \frac{1}{\|w\|} \min_{i=1, \dots, N} |w^T x_i + b|$ subject to $y_i (w^T x_i + b) > 0 \quad \forall i$

Optimization Problem $\min_{w, b} \frac{1}{2} \|w\|^2$ subject to: $y_i (w^T x_i + b) \geq 1 \quad \forall i$. Introducing Lagrange multipliers $\alpha_i \geq 0$, the Lagrangian is: $\mathcal{L}(w, b, \alpha) = \frac{1}{2} \|w\|^2 - \sum_{i=1}^N \alpha_i [y_i (w^T x_i + b) - 1]$

Karush-Kuhn-Tucker (KKT) Conditions: Define the Lagrangian $\mathcal{L}(z, \mu) = F(z) + \mu^T G(z)$. Then, under technical conditions, for each locally optimal \hat{z} , there exists Lagrange multipliers $\hat{\mu} \in \mathbb{R}^m$ such that:

stationarity $\nabla_z \mathcal{L}(\hat{z}, \hat{\mu}) = 0$

Primal Feasibility $G(\hat{z}) \leq 0$

Dual Feasibility $\hat{\mu} \geq 0$

Complementary Slackness $\hat{\mu}_j G_j(\hat{z}_j) = 0$

Dual Problem $\max_{\mu \geq 0} \hat{F}(\mu)$ where $\hat{F}(\mu) = \min_z \mathcal{L}(z, \mu) \mathcal{L}(z, \mu) = F(z) + \mu^T G(z)$

KKT conditions for SVM

1. **From Stationarity** $\hat{w} = \sum_{i=1}^N \hat{\mu}_i y_i x_i, \quad 0 = \sum_{i=1}^N \hat{\mu}_i y_i$
2. **From Dual Feasibility** $\hat{\mu}_i \geq 0$ for $i = 1, \dots, N$

3. **From Complementary Slackness** $\hat{\mu}_i = 0$ or $y_i (\hat{w}^T x_i + \hat{b}) = 1$

4. **The multipliers** $\hat{\mu}$ can be found by the dual problem $\max_{\mu \in \mathbb{R}^N} \sum_{i=1}^N \mu_i - \frac{1}{2} \sum_{i,j} \mu_i \mu_j y_i y_j (x_i^T x_j)$

Dual formulation of SVM

$$\hat{\mu} = \arg \max_{\mu \in \mathbb{R}^N} \sum_{i=1}^N \mu_i - \frac{1}{2} \sum_{i,j} \mu_i \mu_j y_i y_j (x_i^T x_j)$$

Subject to: $\hat{\mu} \geq 0$ and $\sum_{i=1}^N \hat{\mu}_i y_i = 0$

Decision function: $\hat{f}(x) = \text{sgn} \left(\sum_{i=1}^N \hat{\mu}_i y_i x_i^T x + \hat{b} \right)$

Complementary slackness: $\hat{\mu}_i = 0$ or $1 = y_i (\hat{w}^T x_i + \hat{b})$

Kernel Support Vector Machines

$$\hat{\mu} = \arg \max_{\mu \in \mathbb{R}^N} \sum_{i=1}^N \mu_i - \frac{1}{2} \sum_{i,j=1}^N \mu_i \mu_j y_i y_j k(x_i, x_j)$$

Subject to: $\hat{\mu} \geq 0$ and $\sum_{i=1}^N \hat{\mu}_i y_i = 0$

Decision function: $\hat{f}(x) = \text{sgn} \left(\sum_{i=1}^N \hat{\mu}_i y_i k(x_i, x) + \hat{b} \right)$

Only support vectors satisfying $1 = y_i (\hat{w}^T \varphi_i(x) + \hat{b})$ matter for predictions. This is a **sparse kernel method**.

Lecture 4

Decision Trees: Stratify the input space into distinct, non-overlapping regions and assign a chosen, **constant** prediction to each region

Universal approximation theorem In general, any (sufficiently regular) oracle function $f^* : \mathcal{X} \rightarrow \mathbb{R}$ can be approximated with a decision tree with an arbitrary small precision $\epsilon > 0$. i.e. there exists a decision tree T (with some finite depth) such that $|T(x) - f^*(x)| < \epsilon$. Want ϵ to be small \Rightarrow choose larger L

Classification and Regression Trees Suppose that the input space is \mathcal{X} . A **partition** of \mathcal{X} is a collection of subsets $\mathcal{R}_1, \mathcal{R}_2, \dots, \mathcal{R}_J$ such that $\mathcal{R}_i \cap \mathcal{R}_j = \emptyset$ for $i \neq j$ and $\bigcup_{j=1}^J \mathcal{R}_j = \mathcal{X}$

The general decision tree hypothesis space is: $\mathcal{H} = \left\{ f : f(x) = \sum_{j=1}^J a_j \mathbb{1}_{x \in \mathcal{R}_j}, \{\mathcal{R}_j\} \text{ is a partition of } \mathcal{X}, a_j \in \mathbb{Y} \right\}$

where $\mathbb{1}_{x \in \mathcal{R}_j} = \begin{cases} 1, & \text{if } x \in \mathcal{R}_j \\ 0, & \text{otherwise} \end{cases}$

Learning Decision Trees

A decision tree model $f(x) = \sum_{j=1}^J a_j \mathbb{1}_{x \in \mathcal{R}_j}$ depends on both a_j and \mathcal{R}_j . Given $\{\mathcal{R}_j\}$, $\{a_j\}$ are easy to fix:

- **Regression:** we take the **average** label values $a_j = y_j = \frac{\sum_i y_i \mathbb{1}_{x \in \mathcal{R}_j}}{\sum_i \mathbb{1}_{x \in \mathcal{R}_j}}$
- **Classification:** we take the **modal** label values $a_j = \text{mode}\{y_i : x_i \in \mathcal{R}_j\}$

Suppose we are dealing with regression, then we can fix a_j as before and solve the following empirical risk minimization:

$$\min_{\mathcal{R}_j} \frac{1}{2} \sum_{i=1}^N \left(\sum_{j=1}^J a_j \mathbb{1}_{x \in \mathcal{R}_j} - y_i \right)^2. \text{ However this is very hard to solve (NP-hard).}$$

Recursive Binary Splitting Greedy algorithm which essentially repeats the following two steps and grows the tree by adding two leaf nodes at a time: Pick a dimension of the input space, then find the **best** value θ to split this input dimension into two parts and assign new constant values to these new regions.

Loss function for Decision Trees Classification Entropy: $-\sum_{k=1}^K \sum_{j=1}^J p_{jk} \log p_{jk}$ Gini Impurity: $\sum_{k=1}^K \sum_{j=1}^J p_{jk}(1 - p_{jk})$ where p_{jk} is the proportion of samples in \mathcal{R}_j belonging to class k .

Advantages of Decision Trees: 1.Can readily visualize and understand predictions; 2.Implicit feature selection via analyzing contribution of splits to reduction of error/impurity; 3.Robust to data types, supervised learning tasks, and nonlinear relationships.

Disadvantages of Decision Trees: 1.Greedy algorithms may find sub-optimal solutions; 2.Sensitive to data variation and balancing; 3.Prone to overfitting.

Model Ensembling Bagging reduces the variance by aggregating predictions from multiple models trained on different random subsamples of the data:

- **Regression** $\bar{f}(x) = \frac{1}{m} \sum_{j=1}^m f_j(x)$
- **Classification** $\bar{f}(x) = \text{Mode}\{f_j(x) : j = 1, \dots, m\}$

Boosting works by training weak learners sequentially, where each learner focuses on correcting the mistakes of the previous ones. The combined model is a weighted sum of the weak learners: $f(x) = \sum_{t=1}^T \alpha_t f_t(x)$ where α_t are coefficients based on each learner's performance. Boosting helps reduce bias.

Key Ideas of AdaBoost 1. Initialize with uniform weight across all training samples

2. Train a classifier/regressor f_1
3. Identify the samples that f_1 got wrong (classification) or has large errors (regression)
4. Weight these samples more heavily and train f_2 on this reweighted dataset
5. Repeat steps 3-5

AdaBoost Implementation

- **Data:** $\mathcal{D} = \{(x_i, y_i)\}_{i=1}^N$
- Initialize $w_i^{(1)} = \frac{1}{N}$ for all $i = 1, \dots, N$;
- For $j = 1, \dots, m$ do

1. Obtain f_j from: $f_j = \arg \min_{f \in \mathcal{H}} \sum_{i=1}^N w_i^{(j)} \mathbb{1}_{y_i \neq f(x_i)}$
2. Compute combination coefficients: $\delta_j = \frac{\sum_{i=1}^N w_i^{(j)} \mathbb{1}_{y_i \neq f_j(x_i)}}{\sum_{i=1}^N w_i^{(j)}} \quad \alpha_j = \log \left(\frac{1 - \delta_j}{\delta_j} \right)$
3. Update weights: $w_i^{(j+1)} = w_i^{(j)} \exp \left(\alpha_j \mathbb{1}_{y_i \neq f_j(x_i)} \right)$

- Return: $\bar{f}(x) = \text{Sign} \left(\sum_{j=1}^m \alpha_j f_j(x) \right)$

Cross-Validation Cross-validation is essential for tuning model hyperparameters. In k -fold cross-validation, the data is split into k subsets. The model is trained on $k - 1$ subsets and validated on the remaining one. This process is repeated k -times, and the average performance is evaluated. Cross-validation ensures model generalizability.

Neural Network Architecture Example

The architecture of the network consists of: $x = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}$, $x_1 \in \mathbb{R}$, $x_2 \in \mathbb{R}$ **Input layer:** $x_0 = x = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} \in \mathbb{R}^2$

1st hidden layer with 2 neurons. 2nd hidden layer with 3 neurons. Output: Scalar.

1st Hidden Layer $x_1 = W_0 x_0 + b_0$ $W_0 \in \mathbb{R}^{2 \times 2}$, $b_0 \in \mathbb{R}^{2 \times 1}$ $x_1 = \begin{bmatrix} x_{11} \\ x_{12} \end{bmatrix}$, $x_1 \in \mathbb{R}^{2 \times 1}$ $x_1 = \text{ReLU}(x_1)$

2nd Hidden Layer $x_2 = W_1 x_1 + b_1$ $W_1 \in \mathbb{R}^{3 \times 2}$, $b_1 \in \mathbb{R}^{3 \times 1}$ $x_2 = \text{ReLU}(x_2)$, $x_2 \in \mathbb{R}^{3 \times 1}$

Back-propagation Example Computation

The model is defined as: $y(x) = v \delta(w_1 \delta(w_0 x))$, $x, w_0, w_1, v \in \mathbb{R}$, where δ is the identity function, i.e., $\delta(z) = z$.

The loss function is: $L = (y(x) - y)^2$ We want to compute the gradients: $\frac{\partial L}{\partial v}$, $\frac{\partial L}{\partial w_1}$, $\frac{\partial L}{\partial w_0}$

Forward pass: $x_1 = w_0 x$ $x_2 = v w_1 x_1$ $L = (x_2 - y)^2$

Backward pass:

$$p_2 = \frac{\partial L}{\partial x_2} = 2(x_2 - y) \quad \frac{\partial L}{\partial v} = \frac{\partial L}{\partial x_2} \cdot \frac{\partial x_2}{\partial v} = 2(x_2 - y) w_1 x_1 \quad \frac{\partial L}{\partial w_1} = \frac{\partial L}{\partial x_2} \cdot \frac{\partial x_2}{\partial w_1} = 2(x_2 - y) v x_1$$

$$p_1 = \frac{\partial L}{\partial x_1} = \frac{\partial L}{\partial x_2} \cdot \frac{\partial x_2}{\partial x_1} = 2(x_2 - y) v w_1 \quad \frac{\partial L}{\partial w_0} = \frac{\partial L}{\partial x_1} \cdot \frac{\partial x_1}{\partial w_0} = 2(x_2 - y) v w_1 x$$

Lecture 5

Neural Networks The neural network (NN) hypothesis space is quite like linear basis models: $f(x) = \sum_{j=1}^M v_j \phi_j(x)$,

where $\phi_j(x) = \sigma(w_j^T x + b_j)$.

$w_j \in \mathbb{R}^d$: the trainable weights of the hidden layer $b_j \in \mathbb{R}$: the trainable biases of the hidden layer

$v_j \in \mathbb{R}$: the trainable weights of the output layer $\sigma : \mathbb{R} \rightarrow \mathbb{R}$ is the activation function.

Activation Functions

Sigmoid: $\sigma(z) = \frac{1}{1+e^{-z}}$, $[0, 1]$ **Tanh:** $\sigma(z) = \tanh(z)$, $[-1, 1]$

Rectified Linear Unit (ReLU): $\sigma(z) = \max(0, z)$, $[0, \infty]$ **Leaky-ReLU:** $\sigma(z) = \begin{cases} z & \text{if } z \geq 0 \\ \delta z & \text{if } z < 0 \end{cases}$, $[-\infty, \infty]$

Universal Approximation Theorem Any continuous function f^* on a compact domain can be approximated by neural networks to arbitrary precision, provided there are enough neurons (M large enough).

Curse of Dimensionality: Exponential increase in number of neurons required as the dimensionality increases. Baron(1993) showed that for any continuous function $f^* : [0, 1]^d \rightarrow \mathbb{R}$ (with other conditions), there exists a width- M neural network f_M such that $\|f^* - f_M\|^2 \leq \mathcal{O}(M^{-1})$

The Significance of Data-Dependent Feature Maps

Functions behave just like vectors: Each ϕ_j is like a coordinate axis. They play the role of e_j , and the oracle function f^* plays the role of u . There are an **infinite** number of them.

Writing $f^*(x) = \sum_{j=1}^M v_j \phi_j(x)$ is like expanding a vector into its components, but we can't have all components since M is finite. If we get to choose which components to have in the sum **after** seeing some information on f^* , we can usually do much better.

Gradient Descent $w^{(t+1)} = w^{(t)} - \eta \frac{\partial L}{\partial w}$ When η is too small, updates are **slow**; When η is too large, the updates may become **unstable**.

Stochastic Gradient Descent (SGD): $\theta_{k+1} = \theta_k - \eta \frac{1}{B} \sum_{i \in I_B} \nabla \Phi_i(\theta_k)$, this diminishes the probability of sticking at local minima as the main drawback for GD.

Convex function is a function satisfying: $\Phi(\lambda \theta + (1 - \lambda) \theta') \leq \lambda \Phi(\theta) + (1 - \lambda) \Phi(\theta')$ for all $\theta, \theta' \in \mathbb{R}^p$ and $\lambda \in [0, 1]$

Deep Neural Networks (DNNs)

Deep neural networks are an extension of shallow networks. The idea is to stack many hidden layers together

$x_0 = x$

$x_1 = \sigma(W_0 x_0 + b_0)$

$x_T = \sigma(W_{T-1} x_{T-1} + b_{T-1})$

$f(x) = v^T x_T$

where x_0 is the input, x_1, x_2, \dots, x_T are the outputs of hidden layers. W_i and b_i are the weights and biases for each layer, σ is the activation function, $f(x)$ is the final output.

Back-propagation Algorithm

- Initialize $x_0 = x \in \mathbb{R}^d$.
- For $t = 0, 1, \dots, T$: $x_{t+1} = g_t(x_t, W_t) = \sigma(W_t x_t)$
- Set $p_{T+1} = \nabla_x L(x_{T+1}, y)$.
- For $t = T, T-1, \dots, 1$:
 - $\nabla W_t \Phi = p_{t+1}^T \nabla_W g_t(x_t, W_t)$
 - $p_t = [\nabla_x g_t(x_t, W_t)]^T p_{t+1}$
- Return $\{\nabla W_t \Phi : t = 0, \dots, T\}$.

Unsupervised Learning

Lecture 6

Principal Component Analysis

- Find the direction that captures the most variance, where the order in magnitude of the normalized sample covariance matrix S corresponds to the order of eigenvalues. λ
- Find the direction that minimizes projection error by partitioning the set of eigenvectors and choosing the exclusion set ($m+1$ to d) to be the eigenvector subset corresponding to the smallest ($d - m - 1$) eigenvalues.

PCA Algorithm Implementation

- **Data:** $\mathcal{D} = \{x_i\}_{i=1}^N$, $x_i \in \mathbb{R}^d$ for all i .
- **Hyper-parameters:** m (reduction dimension).
- Compute sample covariance matrix $S = \frac{1}{N} \sum_{i=1}^N x_i x_i^T$.
- Compute the first m eigenvectors $\{u_1, \dots, u_m\}$ and eigenvalues $\{\lambda_1, \dots, \lambda_m\}$.
- Form $d \times m$ matrix U_m whose j^{th} column is u_j .
- Compute $Z_m = X U_m$.
- Return Principal component scores Z_m , Eigenvalues, and eigenvectors λ_j, u_j for $j = 1, \dots, m$.

PCA in Feature Space

- **Data:** $\mathcal{D} = \{x_i\}_{i=1}^N$, $x_i \in \mathbb{R}^d$ for all i
- **Hyper-parameters:** m (reduction dimension), ϕ (feature maps);
- Compute design matrix $\Phi_{ij} = \phi_j(x_i)$;
- Center design matrix $\Phi_{ij} \leftarrow \Phi_{ij} - \frac{1}{N} \sum_{i=1}^N \Phi_{ij}$;
- Compute sample covariance matrix $S_\phi = \frac{1}{N} \Phi^T \Phi$;
- Compute the first m eigenvectors $\{u_1, \dots, u_m\}$ and eigenvalues $\{\lambda_1, \dots, \lambda_m\}$ of S_ϕ ;
- Form $d \times m$ matrix U_m whose j^{th} column is u_j ;
- Compute $Z_m = \Phi U_m$;
- **Return:** Principal component score Z_m , Eigenvalues and eigenvectors λ_j, u_j for $j = 1, \dots, m$.

PCA whitening transform Principal component scores are given by $Z = X U$ where X is the original features and U is the matrix of eigenvectors. The transformation $X' = X U \Lambda^{-\frac{1}{2}}$, where Λ is the matrix of eigenvalues makes $\text{cov}(X') = I$.

Useful derivatives

$$\frac{\partial}{\partial x} (x^T A x) = 2 A x \quad \frac{\partial}{\partial x} (A^T x) = A$$

PCA Example

Flow: Center data -> compute sample covariance matrix -> compute eigenvectors of largest eigenvalues
Consider a simple dataset with four 2D points:

$$\begin{aligned}x_1 &= (1, 1) \\x_2 &= (-1, -1) \\x_3 &= (\epsilon, -\epsilon) \\x_4 &= (-\epsilon, \epsilon)\end{aligned}$$

where ϵ is a small positive number.

Step 1: Compute the sample covariance matrix S

$$\begin{aligned}S &= \frac{1}{4} \left[(1, 1)^T (1, 1) + (-1, -1)^T (-1, -1) + (\epsilon, -\epsilon)^T (\epsilon, -\epsilon) + (-\epsilon, \epsilon)^T (-\epsilon, \epsilon) \right] \\&= \frac{1}{4} \left[\begin{pmatrix} 1 & 1 \\ 1 & 1 \end{pmatrix} + \begin{pmatrix} 1 & 1 \\ 1 & 1 \end{pmatrix} + \begin{pmatrix} \epsilon^2 & -\epsilon^2 \\ -\epsilon^2 & \epsilon^2 \end{pmatrix} + \begin{pmatrix} \epsilon^2 & -\epsilon^2 \\ -\epsilon^2 & \epsilon^2 \end{pmatrix} \right] \\&= \begin{pmatrix} 1 + \epsilon^2 & 1 - \epsilon^2 \\ 1 - \epsilon^2 & 1 + \epsilon^2 \end{pmatrix}\end{aligned}$$

Step 2: Find eigenvalues and eigenvectors

$$\begin{aligned}\det(S - \lambda I) &= (1 + \epsilon^2 - \lambda)^2 - (1 - \epsilon^2)^2 = 0 \\ \lambda_1 &= 2, \lambda_2 = 2\epsilon^2\end{aligned}$$

For $\lambda_1 = 2$:

$$\begin{pmatrix} 1 + \epsilon^2 - 2 & 1 - \epsilon^2 \\ 1 - \epsilon^2 & 1 + \epsilon^2 - 2 \end{pmatrix} \begin{pmatrix} u_1 \\ u_2 \end{pmatrix} = \begin{pmatrix} 0 \\ 0 \end{pmatrix}$$

Solving this, we get the eigenvector: $u = \begin{pmatrix} 1/\sqrt{2} \\ 1/\sqrt{2} \end{pmatrix}$

Step 3: Project data onto the first principal component

$$\begin{aligned}z_1 &= (1, 1) \begin{pmatrix} 1/\sqrt{2} \\ 1/\sqrt{2} \end{pmatrix} = \sqrt{2} \\ z_2 &= (-1, -1) \begin{pmatrix} 1/\sqrt{2} \\ 1/\sqrt{2} \end{pmatrix} = -\sqrt{2} \\ z_3 &= (\epsilon, -\epsilon) \begin{pmatrix} 1/\sqrt{2} \\ 1/\sqrt{2} \end{pmatrix} = 0 \\ z_4 &= (-\epsilon, \epsilon) \begin{pmatrix} 1/\sqrt{2} \\ 1/\sqrt{2} \end{pmatrix} = 0\end{aligned}$$